

Data Structures and Algorithms


Week 6

Graph Algorithms


Basics

- A *graph* is a set V of *vertices* and a collection E of *edges*, $G = (V, E)$
- An edge connecting vertices (or nodes) u and v is denoted (u, v) .
- An edge can be *directed* or *undirected*.
- Directed graph vs. undirected graph:

(a) Undirected graph




(b) Directed graph



Graph Algorithms

Basics


- Two vertices u and v are said to be *adjacent* if there is an edge (u, v) .
- An edge is said to be *incident* to a vertex if the vertex is one of the edge's endpoints.
- *Weighted* graph: An information (usually called weight) is associated with edges



Graph Algorithms

Basics


- Outgoing edge vs. incoming edge
 - Degree, in-degree, and out-degree of a node
-
- The outgoing edges of vertex g are (g, e) , (g, h) .
 - The incoming edges of vertex g are (a, g) , (b, g) , (c, g) .
 - The degree of vertex g , $\deg(g) = 5$.
 - The in-degree of vertex g , $\text{indeg}(g) = 3$.
 - The out-degree of vertex g , $\text{outdeg}(g) = 2$.



Graph Algorithms

Basics

- Parallel edges and self-loops




Graph Algorithms


Basics

- Path, cycle, simple path, simple cycle, directed path, directed cycle*

(a) Undirected graph



(b) Directed graph




Graph Algorithms


Basics

- *Connected graph and strongly connected graph*

(a) Connected graph



(b) Strongly connected graph




Graph Algorithms


Basics

- *Subgraph and spanning subgraph*

(a) A subgraph



(b) A spanning subgraph




Graph Algorithms


Basics

- *Forest and tree*

(a) A forest



(b) A tree



- A *spanning tree* of a graph is a spanning subgraph that is a tree

Graph Algorithms

Basics

- Graph properties
 - If a graph $G = (V, E)$ has m edges, then
$$\sum_{v \in V} \deg(v) = 2m$$
 - If $G = (V, E)$ is a directed graph with m edges, then
$$\sum_{v \in V} \text{in deg}(v) = \sum_{v \in V} \text{out deg}(v) = m$$
- Let G be a simple graph with n vertices and m edges.
 - If G is undirected, then $m \leq \frac{n(n-1)}{2}$
 - If G is directed, then $m \leq n(n-1)$.

Graph Algorithms

Basics

- Graph properties (continued)
 - Let G be an undirected graph with n vertices and m edges:
 - If G is connected, then $m \geq n - 1$
 - If G is a tree, then $m = n - 1$
 - If G is a forest, then $m \leq n - 1$

Graph Algorithms

Graph ADT

- Operations
 - numVertices()
 - vertices()
 - numEdges()
 - edges()
 - getEdge(u , v)
 - endVertices(e)
 - opposite(v , e)

Graph Algorithms

Graph ADT

- Operations (continued)
 - `outDegree(v)`
 - `indegree(v)`
 - `outgoingEdges(v)`
 - `incomingEdges(v)`
 - `insertVertex(x)`
 - `insertEdge(u, v, x)`
 - `removeVertex(v)`
 - `removeEdge(e)`

Graph Algorithms

Data Structures for Graphs

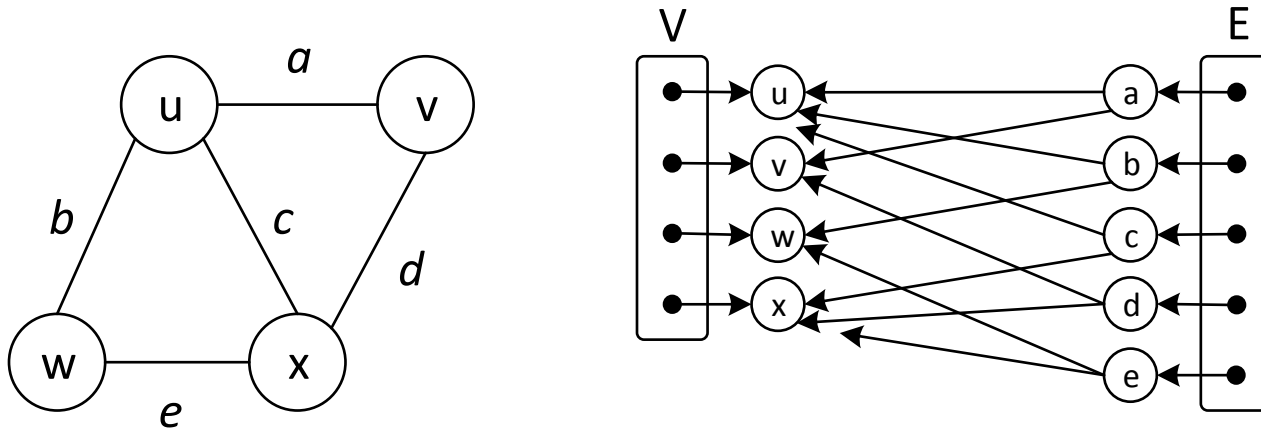
- Edge list, adjacency list, adjacency map, adjacency matrix

Method	Edge List	Adj. List	Adj. Map	Adj. Matrix
vertices()	$O(n)$	$O(n)$	$O(n)$	$O(n)$
edges()	$O(m)$	$O(m)$	$O(m)$	$O(m)$
getEdge(u, v)	$O(m)$	$O(\min(d_u, d_v))$	$O(1)$ exp.	$O(1)$
outDegree(v) inDegree(v)	$O(m)$	$O(1)$	$O(1)$	$O(n)$
outgoingEdges(v) incomingEdges(v)	$O(m)$	$O(d_v)$	$O(d_v)$	$O(n)$
insertVertex(x)	$O(1)$	$O(1)$	$O(1)$	$O(n^2)$
removeVertex(v)	$O(m)$	$O(d_v)$	$O(d_v)$	$O(n^2)$
insertEdge(u, v, x)	$O(1)$	$O(1)$	$O(1)$ exp.	$O(1)$
remove Edge(e)	$O(1)$	$O(1)$	$O(1)$ exp.	$O(1)$

Graph Algorithms

Data Structures for Graphs

- Edge list




- V is a list of vertices and E is a list of edges. Both can be implemented using doubly linked lists.

Graph Algorithms

Data Structures for Graphs

- Adjacency list

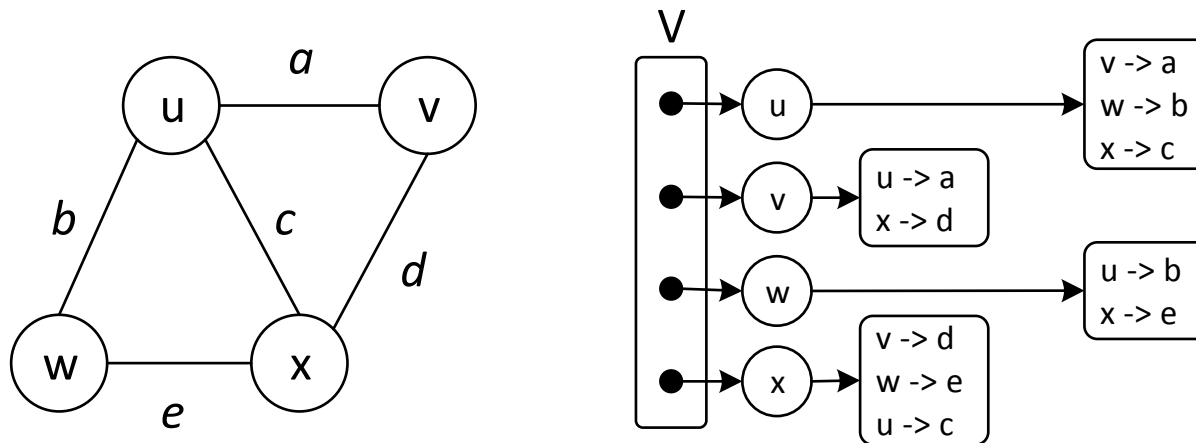


- V is a list of vertices.
- Each vertex v has a reference to a separate collection of edges that are incident to v .
- The collection is called *incidence collection*.

Graph Algorithms

Data Structures for Graphs

- Adjacency map




- Incidence collection of V is implemented as a map.
- Suppose edge $a = (u, v)$ is in the incidence collection of u . Then, $\langle v, a \rangle$ pair is stored in the map, where v is a key and a is the corresponding value.

Graph Algorithms

Data Structures for Graphs

- Adjacency matrix



- $n \times n$ matrix.
- Vertices are encoded to integers and these integers are used as indexes.
- The entry corresponding to vertices u and v stores an edge (u, v) .

Graph Algorithms

Graph Traversals

- A *graph traversal* is a systematic procedure for visiting (and processing) all vertices in the graph.
- We say a traversal is efficient if its running time is proportional to the number of vertices and edges in the graph.
- Applications (for directed graph):
 - Find a direct path from vertex u to vertex v .
 - Find all vertices of G that are reachable from a given vertex s .
 - Determine whether G is acyclic.
 - Determine whether G is strongly connected.

Graph Algorithms

Graph Traversals

- Applications (for undirected graph):
 - Find a path from vertex u to vertex v .
 - Given a start vertex s , find a path with the minimum number of edges from s to every other vertex.
 - Test whether G is connected.
 - Find a spanning tree of G .
 - Identify a cycle in G .
- Will discuss *depth-first search (DFS)* and *bread-first search (BFS)*.

Graph Algorithms

DFS

- Pseudocode

Algorithm DFS (G, u)

Input: A graph G and a vertex u of G

Output: A collection of vertices reachable from u , with
their discovery edges

Mark u as visited

for each of u 's outgoing edges, $e = (u, v)$ do

 if v has not been visited then

 Record edge e as the discovery edge for vertex v

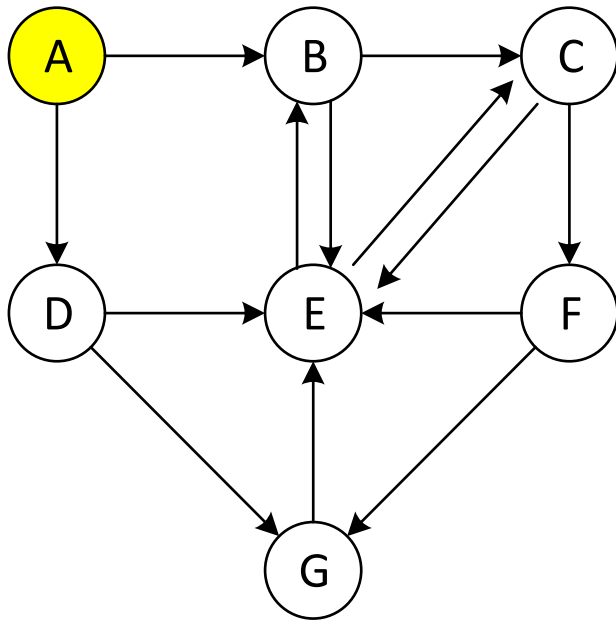
 Recursively call DFS(G, v)

Graph Algorithms

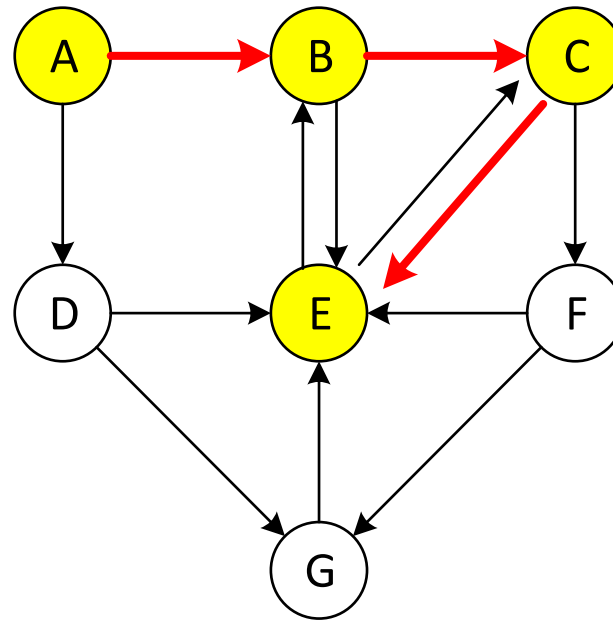
DFS

- Illustration (on a directed graph)

A directed graph
Start at vertex A



A → B → C → E
Backtrack to C



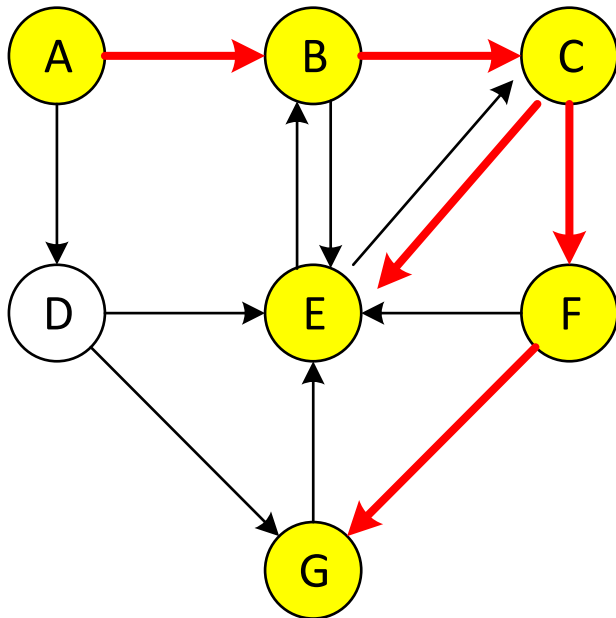
Graph Algorithms

DFS

- Illustration (on a directed graph)


C -> F -> G

Backtrack to F -> C -> B -> A



A -> D

Finished



Graph Algorithms

DFS

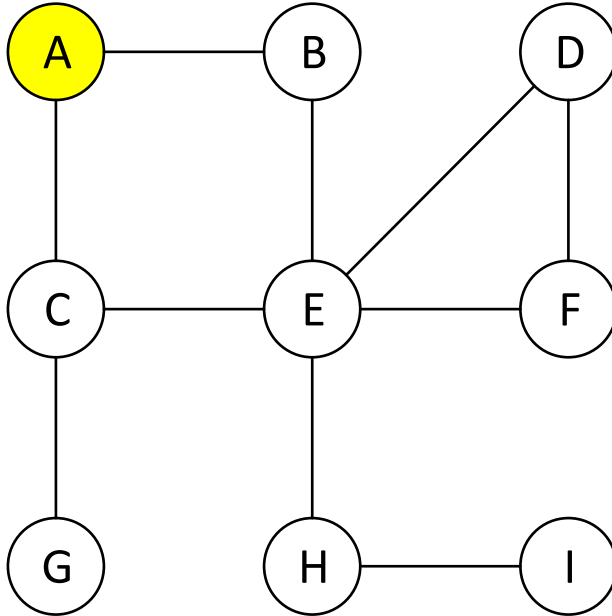
- Illustration (on a directed graph)
 - Classification of edges:
 - *Back edges*: A back edge connects a vertex to its ancestor in the *DFS* tree. They are labeled *B*.
 - *Forward edges*: A forward edge connects a vertex to its descendant in the *DFS* tree. They are labeled *F*.
 - *Cross edge*: A cross edge connects a vertex to a vertex that is neither its ancestor nor its descendant. They are labeled *C*.

Graph Algorithms


DFS

- Illustration (on an undirected graph)

An undirected graph
Start at vertex A



A -> B -> E -> C -> G
Backtrack to C -> E



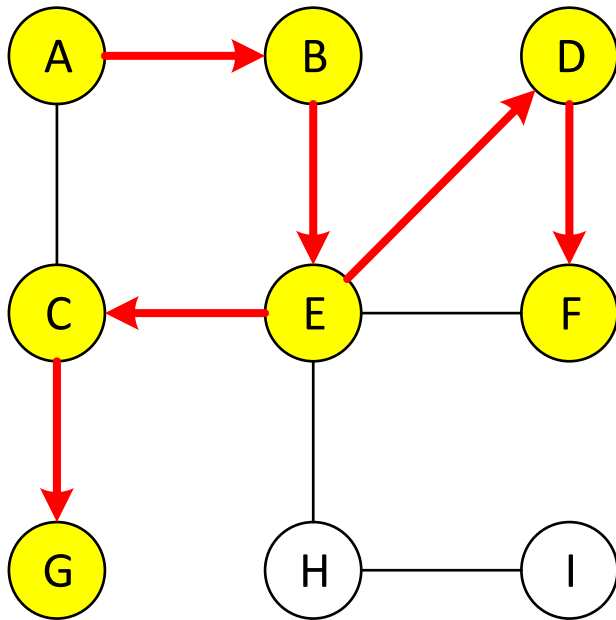
Graph Algorithms

DFS

- Illustration (on an undirected graph)

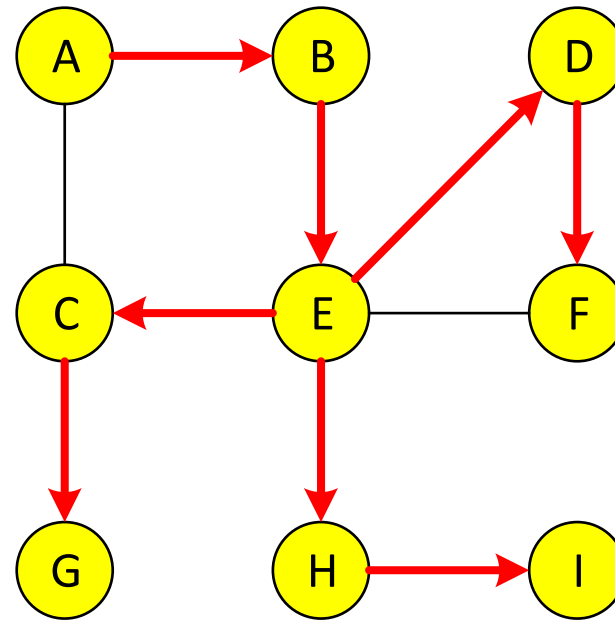
E -> D -> F

Backtrack to D -> E



E -> H -> I

Finished



Graph Algorithms


DFS

- DFS properties:
 - A *DFS* on an undirected graph G starting at a vertex s visits all vertices in the connected component of s , and the discovery edges form a *spanning tree* of the connected component of s .
 - A *DFS* on a directed graph G starting at a vertex s visits all vertices reachable from s , and the *DFS* tree contains the directed paths from s to every vertex reachable from s .
- Running time: $O(n_s + m_s)$, here n_s is the number of vertices reachable from s and m_s is the number of edges that are incident to those vertices


Graph Algorithms

BFS

- Outline
 - Start at the starting vertex s
 - Visit all vertices that are “one-edge away” from s
 - Visit all vertices that are “two-edge away” from s
 - and so on.
- Illustration



Explore vertices that are one-edge away from A.




Graph Algorithms


BFS

- Illustration (continued)


Explore vertices that are two-edge away from A.



Explore vertices that are three-edge away from A.



Explore vertices that are four-edge away from A. Finished



Graph Algorithms


BFS

- BFS properties:
 - The traversal visits all vertices reachable from s .
 - For each vertex v at level i , the path in the *BFS* tree from s to v has i edges, and any other path from s to v in G has at least i edges.
 - If (u, v) is an edge that is not in the *BFS* tree, the level number of v is at most 1 greater than the level number of u .
- Running time: $O(n + m)$

Graph Algorithms

Weighted Graph

- Each edge e is associated with a numeric label called *weight*, denoted $w(e)$.
- Example



Graph Algorithms

Shortest Paths

- Let G be a weighted graph.
 - The *length* of a path P is the sum of the weights of all edges on P . Let $P = \langle (v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k) \rangle$. Then, the length of P , denoted $w(P)$, is defined as:


$$w(P) = \sum_{i=0}^{k-1} w(v_i, v_{i+1})$$

- The *distance* from a vertex u to a vertex v in G , $d(u, v)$, is the length of a minimum-length path from u to v , if such path exists. The minimum-length path is referred to as *shortest path*.
 - $d(u, v) = \infty$, if there is no path from u to v in G .

Graph Algorithms

Shortest Paths

- Weights can be negative numbers. Then, a graph may have a *negative-weight cycle*:



- If a graph has a negative-weight cycle, a shortest path is not well defined.

Graph Algorithms

Dijkstra's Algorithm

- A well-known single-source shortest path algorithm on a directed or undirected graph G without negative weights.
- Finds shortest paths from a source vertex to every other vertex in G .
- A greedy algorithm.
- Edge relaxation
 - $D[v]$ is the length of the best path from s to v we have found so far.
 - Initially $D[s] = 0$ and $D[v] = \infty$ for all other vertexes.
 - During the execution of the algorithm, $D[v]$ is updated iteratively and becomes a shortest-path length from s to v .

Graph Algorithms


Dijkstra's Algorithm

- Edge relaxation (continued)


if $D[u] + w(u, v) < D[v]$ then

$$D[v] = D[u] + w(u, v)$$


before relaxation



after relaxation



not relaxed



Graph Algorithms

Dijkstra's Algorithm

- Pseudocode

Algorithm ShortestPath(G, s):

Input: A directed or undirected graph G with nonnegative weights, and a distinguished vertex s of G

Output: The length of a shortest path from s to v for every vertex v of G

Initialize $D[s] = 0$ and $D[v] = \infty$ for each vertex $v \neq s$

Let a priority queue Q contains all vertices of G using D labels as keys

while Q is not empty do

$u = Q.\text{removeMin}()$ // vertex with the smallest $D[u]$ is pulled into “cloud”

 for each edge (u, v) such that v is in Q do

 // perform relaxation

 if $D[u] + w(u, v) < D[v]$ then

$D[v] = D[u] + w(u, v)$


 Change the key of vertex v in Q to $D[v]$

return the label $D[v]$ of each vertex v


Graph Algorithms

Dijkstra's Algorithm


- Illustration



(a) Initially, all vertices are in Q, C is empty, $D[s] = 0$, $D[v] = \infty$ for all other vertices.



(b) s comes into C, edges (s, a) and (s, b) are relaxed.




(c) a comes into C, edges (a, b), (a, c), and (a, d) are relaxed.


Graph Algorithms

Dijkstra's Algorithm


- Illustration (continued)



(d) b comes into C, edge (b, e) is relaxed.



(e) d comes into C, edge (d, c) is relaxed.




(f) c comes into C. No edge relaxation needed.

Graph Algorithms

Dijkstra's Algorithm

- Illustration (continued)



Running time: $O((n + m) \log n)$

(g) e comes into C. No edge relaxation needed. Finished.

Graph Algorithms

Minimum Spanning Trees

- Given a tree T in an undirected, weighted graph G , the weight of T , $w(T)$, is defined as follows:


$$w(T) = \sum_{(u,v) \text{ in } T} w(u,v)$$

- A *minimum spanning tree* of an undirected, weighted graph G is a spanning tree with the minimum weight.
- Minimum spanning tree problem: Find such a tree in G .
- Will discuss two algorithms, Prim-Jarnik algorithm and Kruskal's algorithm, both of which are greedy algorithms.
- We assume that a graph G is undirected, weighted, connected, and simple.

Graph Algorithms

Minimum Spanning Trees

- Bridge edge and minimum-weight (bridge) edge
- Suppose G is partitioned into mutually exclusive V_1 and V_2 .
- Bridge edge: one end in V_1 and the other in V_2 .
- Minimum-weight edge: a bridge with the smallest weight



Graph Algorithms

Prim-Jarnik Algorithm

- Outline
 - Begins at some “root” vertex s .
 - Keeps a set of vertices C , called “cloud.”
 - Initially, C has only s .
 - In each iteration, we find a minimum-weight edge connecting a vertex u in the cloud of C and a vertex v that is outside the cloud.
 - Then, the vertex v is pulled into C
 - This process is repeated until a spanning tree is formed.

Graph Algorithms


Prim-Jarnik Algorithm

- Outline (continued)
 - Each vertex v has a label $D[v]$, which stores the weight of the minimum observed edge connecting v to the cloud C .
 - Vertices that are not in C are stored in a priority queue, where $D[v]$ is used as a key in the queue.
 - If we choose a vertex in the priority queue with the minimum $D[v]$, then it is a minimum-weight edge.
- [Pseudocode](#)


Graph Algorithms

Prim-Jarnik Algorithm

- Illustration



(a) Initial tree




(b) (P,J) is minimum-weight edge.


Graph Algorithms

Prim-Jarnik Algorithm

- Illustration (continued)



(c) (J,W) is minimum-weight edge.




(d) (J,B) is minimum-weight edge.


Graph Algorithms

Prim-Jarnik Algorithm

- Illustration (continued)



(e) (W,O) is minimum-weight edge.




(f) (O,D) is minimum-weight edge.


Graph Algorithms

Prim-Jarnik Algorithm

- Illustration (continued)



(g) (W,M) is minimum-weight edge.




(h) (D,L) is minimum-weight edge.


Graph Algorithms

Prim-Jarnik Algorithm

- Illustration (continued)



(i) (L,S) is minimum-weight edge.



(j) Finished. The thick red edges form a minimum spanning tree T .

- Running time: $O((n + m) \log n)$

Graph Algorithms


Kruskal's Algorithm

- In the Prim-Jarnik's algorithm, there is always a single tree.
- In the Kruskal's algorithm, there are multiple trees, all of which are eventually merged into an MST.
- Outline: Initially, a spanning tree T is empty and each vertex is a “cluster” on its own.
 - Step 1: Find an edge e with the smallest weight.
 - Step 2: If two endpoints of e belong to different clusters, merge those two clusters.
 - Step 3: Include e in T .
 - Step 4: Stop if all vertices are included by T . Otherwise, return to Step 1 and repeat.


Graph Algorithms

Kruskal's Algorithm

- [Pseudocode.](#)
- Illustration



(a) Initial tree. Each vertex is its own cluster. $w(J,P)$ is the smallest.




(b) $w(J,W)$ is the next smallest.


Graph Algorithms

Kruskal's Algorithm

- Illustration (continued)



(c) $w(B,J)$ is the next smallest.




(d) $w(L,S)$ is the next smallest.


Graph Algorithms

Kruskal's Algorithm

- Illustration (continued)



(e) $w(O,W)$ is the next smallest.




(f) $w(J,O)$ is the next smallest. But, they are in the same cluster. $w(O,P)$ the same. $w(D,O)$ is the next smallest.


Graph Algorithms

Kruskal's Algorithm

- Illustration (continued)



(g) $w(B,O)$ is the next smallest. But, they are in the same cluster. $w(M,W)$ is the next smallest.




(h) $w(J,M)$ is the next smallest. But, they are in the same cluster. $w(D,M)$ the same. $w(D,L)$ is the next smallest.

Graph Algorithms

Kruskal's Algorithm

- Illustration (continued)



Running time: $O(m \log n)$

(i) Finished. Thick red edges form a minimum spanning tree.

Decision Problem

- Decision problem: A decision problem **P** is a set of questions each of which has a yes or no answer.
- Example: A decision problem **P_{sq}**: Determine whether an arbitrary number is a perfect square or not. This problem consists of the following questions:

p₀: Is 0 a perfect square?

p₁: Is 1 a perfect square?

...

Here, **p_i** is also called an instance of **P**.

Decision Problem


- A solution to a decision problem is an algorithm that determines the answer to every question $\mathbf{p}_i \in \mathbf{P}$.
- An algorithm that solves a decision problem should be
 - *complete* – it produces an answer, either positive or negative, to each question in the problem domain
 - *mechanistic* – it consists of a finite sequence of instructions each of which can be carried out without requiring insight, ingenuity, or guesswork
 - *deterministic* – when presented with identical input, it always produces the same result.

Decision Problem

- Decision problems:
 - *Unsolvable* (or *undecidable*)
 - *Solvable*:
 - *Tractable*: A decision problem is said to be tractable if there is at least one polynomially bounded algorithm that solves the problem. Such an algorithm is called an *efficient* algorithm.
 - *Intractable*: A decision problem is said to be intractable if there is no polynomially bounded algorithm (or no efficient algorithm) that solves the problem

Reducibility

- A decision problem \mathbf{P} is Turing reducible to a problem \mathbf{P}' if there is a Turing machine that takes any problem $\mathbf{p}_i \in \mathbf{P}$ as input and produces an associated problem $\mathbf{p}'_i \in \mathbf{P}'$ where the answer to the original problem \mathbf{p}_i can be obtained from the answer to \mathbf{p}'_i .



P and NP

- A language L is decidable in polynomial time if there is a standard (or deterministic) Turing machine M that accepts L in polynomial time, or $O(n^r)$, where r is a natural number independent of n .
- The family of languages decidable in polynomial time is denoted **P** .

P and NP

- Nondeterministic computation:
 - A deterministic machine solves a decision problem by generating a solution.
 - A nondeterministic machine needs only determine if one of possibilities is a solution.
- A language L is said to be accepted in nondeterministic polynomial time if there is a nondeterministic Turing machine that accepts L in polynomial time, or $O(n^r)$, where r is a natural number independent of n .

P and NP

- The family of languages accepted in nondeterministic polynomial time is denoted **NP**.
- Another definition: A problem is in **NP** if it is “verifiable” in polynomial time.
- What “verifiable” means is that given a possible solution (which is also called **certificate**) we can verify whether it is a solution or not in polynomial time.

P and NP

- **$P = NP$?**
- Unsolved question.
- Since every deterministic machine is also nondeterministic, **$P \subseteq NP$.**
- But it was never proved that **$NP \subseteq P$.** (If this is proved, then that proves **$P = NP$.**)

P and NP

- If Q is reducible to L in polynomial time and $L \in \mathbf{P}$, then $Q \in \mathbf{P}$.
- A language L is called ***NP-hard*** if for every $Q \in \mathbf{NP}$ Q is reducible to L in polynomial time.
- An ***NP-hard*** language that is also in \mathbf{NP} is called ***NP-complete***.
- If there is an NP-complete language that is also in \mathbf{P} , then $\mathbf{P} = \mathbf{NP}$.

P and NP

- Two examples of NP-complete problems: Hamiltonian cycle problem and traveling salesman problem.

Hamiltonian Cycle Problem

- A Hamiltonian cycle of an undirected graph $G = (V, E)$ is a simple cycle that contains each vertex in V .
- Note: A cycle is simple if a node, except the first node, is visited only once.
- A graph that contains a Hamiltonian cycle is called “Hamiltonian.”
- ***Hamiltonian Cycle Problem:*** Does a graph G have a Hamiltonian cycle?

Hamiltonian Cycle Problem


- It can be shown that the Hamiltonian cycle problem can be decidable by a Turing machine in *exponential* time, but not in *polynomial* time. This means Hamiltonian cycle problem is not in **P**.
- But, it is decidable in nondeterministic polynomial time.
- Given a cycle in a graph, we can determine whether it is Hamiltonian cycle or not in polynomial time.
- So, Hamiltonian cycle problem is in **NP**.
- In fact it is an **NP-complete** problem.

Traveling Salesman Problem

- Given a complete, non-negative weighted graph, find a Hamiltonian cycle of minimum weight.
- This problem is ***NP-complete***.
- Will briefly discuss three approximate algorithms.

Traveling Salesman Problem

- Consider the following graph:



minimum weight cycle = $1 \rightarrow 2 \rightarrow 4 \rightarrow 3 \rightarrow 1$.

total weight = $30 + 35 + 22 + 25 = 112$

Traveling Salesman Problem

- Nearest-neighbor strategy

NEAREST-TSP (G, f) /* f is a cost function, or a weight function */

select an arbitrary vertex s ;

$v = s$; $Q = \{v\}$; $S = G.V - Q$; $C = \phi$;

while $S \neq \phi$

select an edge (v, w) of minimum weight, where $w \in S$;

$C = C \cup \{(v, w)\}$;

$Q = Q \cup \{w\}$;

$S = S - \{w\}$;

$v = w$;


$C = C \cup \{(v, s)\}$;

return C ;

Running time: $O(V^2)$

Traveling Salesman Problem

- Nearest-neighbor strategy



Starting at vertex 1: (1, 3), (3, 2), (2, 4), (4, 1)

Total weight = $25 + 20 + 35 + 45 = 125$

Traveling Salesman Problem

- Shortest-link strategy

SHORTEST-LINK-TSP (G, f)

$R = G.E;$

$C = \phi;$

while $R \neq \phi$

 choose the shortest edge (v, w) from R ;

$R = R - \{(v, w)\};$

if (v, w) does not make a cycle with edges in C and (v, w) would
 not be the third edge in C incident on v or w

then

$C = C + \{(v, w)\};$


Running time: $O(E \log V)$

add the edge connecting the end points of the path in C ;

return C ;

Traveling Salesman Problem

- Shortest-link strategy



Edges added: (2, 3), (3, 4), (2, 1), (1, 4)

Total weight = $20 + 22 + 30 + 45 = 117$

Traveling Salesman Problem

- In general, we cannot establish a bound on how much the weight of an approximate algorithm differ from the weight of a minimum tour.
- If we assume the triangle inequality holds on distances among vertices, we can develop an approximate algorithm that has an upper bound on the weight.
- Triangle inequality:
$$f(u, v) \leq f(u, w) + f(w, v), \text{ for all } u, v, w \in G.V.$$
- Euclidean distance has the triangle inequality property.

Traveling Salesman Problem


- The following approximate algorithm has an upper bound on the weight: total weight of a cycle is no more than the twice that of the minimum spanning tree's weight

APPROX-TSP-TOUR (G, f)

```
select a vertex  $r \in G.V$  to be the root;  
compute MST  $T$  from  $r$  using MST-PRIM( $G, f, r$ );  
let  $H$  be a list of vertices, ordered according to when they are  
    first visited in a preorder tree walk of  $T$ ;  
return  $H$ 
```


Traveling Salesman Problem

- Example: Given the following complete graph (There are edges from each node to all other nodes though edges are not shown in the graph below).




Traveling Salesman Problem

- A minimum spanning tree T (a is the root)



Traveling Salesman Problem


- A minimum spanning tree T (a is the root)



$a \rightarrow b \rightarrow c \rightarrow h \rightarrow d \rightarrow e \rightarrow f \rightarrow g \rightarrow a$

Traveling Salesman Problem


- H returned by APPROX-TSP-TOUR is



total weight = approx. 19.074

Traveling Salesman Problem

- An optimal tour (or Hamiltonian cycle with minimum weight)



total weight = approx. 14.715

References

- M.T. Goodrich, R. Tamassia, and M.H. Goldwasser, “Data Structures and Algorithms in Java,” Sixth Edition, Wiley, 2014.
- T.A. Sudkamp, “Languages and Machines,” 1988, Addison Wesley.
- T.H. Cormen, C.E. Leiserson, R.L. Rivest, C. Stein, “Introduction to Algorithms,” 3rd Ed., 2009, MIT Press.